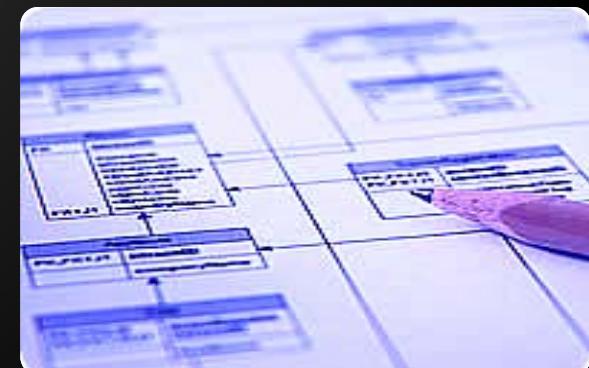
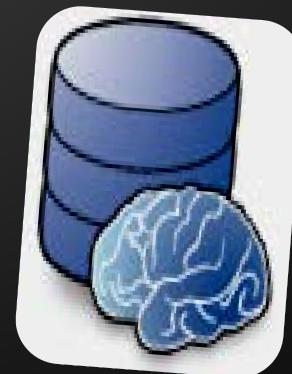




# Database Systems

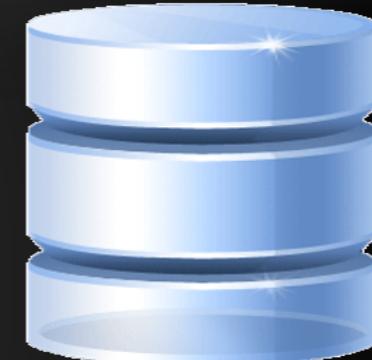
DBMS Fundamental Concepts

---



# Table of Contents

1. Database Models
2. Relational Database Model
3. DBMS & RDBMS Systems
4. Tables, Relationships,  
Multiplicity, E/R Diagrams
5. Normalization
6. Constraints
7. Indices
8. The SQL language



# Table of Contents (2)

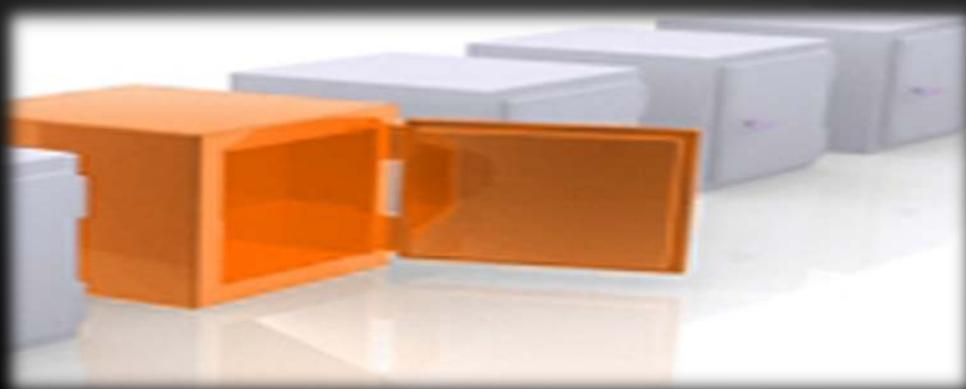
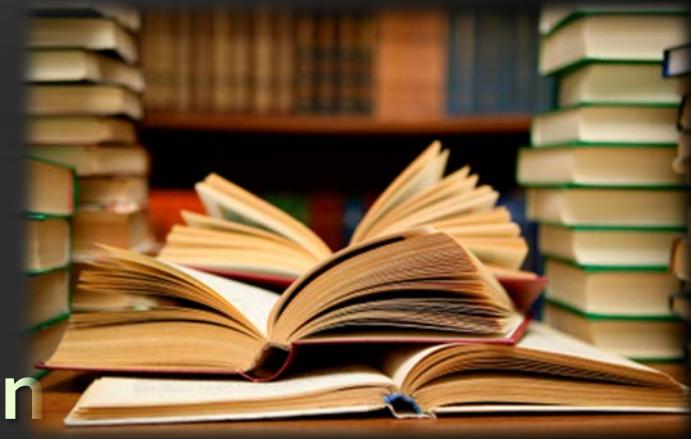
9. Stored Procedures

10. Views

11. Triggers

12. Transactions and Isolation

13. NoSQL Databases



# RDBMS Systems

Relational Databases, Database Servers and RDBMS



# Relational Databases

- ◆ Database models
  - ◆ Hierarchical (tree)
  - ◆ Network / graph
  - ◆ Relational (table)
  - ◆ Object-oriented
- ◆ Relational databases
  - ◆ Represent a bunch of tables together with the relationships between them
  - ◆ Rely on a strong mathematical foundation: the relational algebra



# Relational Database Management System (RDBMS)

- ◆ Relational Database Management Systems (RDBMS) manage data stored in tables
- ◆ RDBMS systems typically implement
  - Creating / altering / deleting tables and relationships between them (database schema)
  - Adding, changing, deleting, searching and retrieving of data stored in the tables
  - Support for the SQL language
  - Transaction management (optional)

# RDBMS Systems

- ◆ RDBMS systems are also known as:
  - ◆ Database management servers
  - ◆ Or just database servers
- ◆ Popular RDBMS servers:
  - ◆ Microsoft SQL Server
  - ◆ Oracle Database
  - ◆ MySQL
  - ◆ IBM DB2
  - ◆ PostgreSQL
  - ◆ SQLite





# Tables and Relationships

Database Tables, Relationships, Multiplicity

- ◆ Database tables consist of data, arranged in rows and columns
  - ◆ For example (table Persons):

<b>Id</b>	<b>First Name</b>	<b>Last Name</b>	<b>Employer</b>
1	Svetlin	Nakov	Telerik
2	Stephen	Forte	Telerik
3	Steve	Jobs	Apple

- ◆ All rows have the same structure
- ◆ Columns have name and type (number, string, date, image, or other)

- ◆ The schema of a table is an ordered sequence of column specifications (name and type)
- ◆ For example the Persons table has the following schema:

```
Persons (
    Id: number,
    FirstName: string,
    LastName: string,
    Employer: string
)
```

# Primary Key

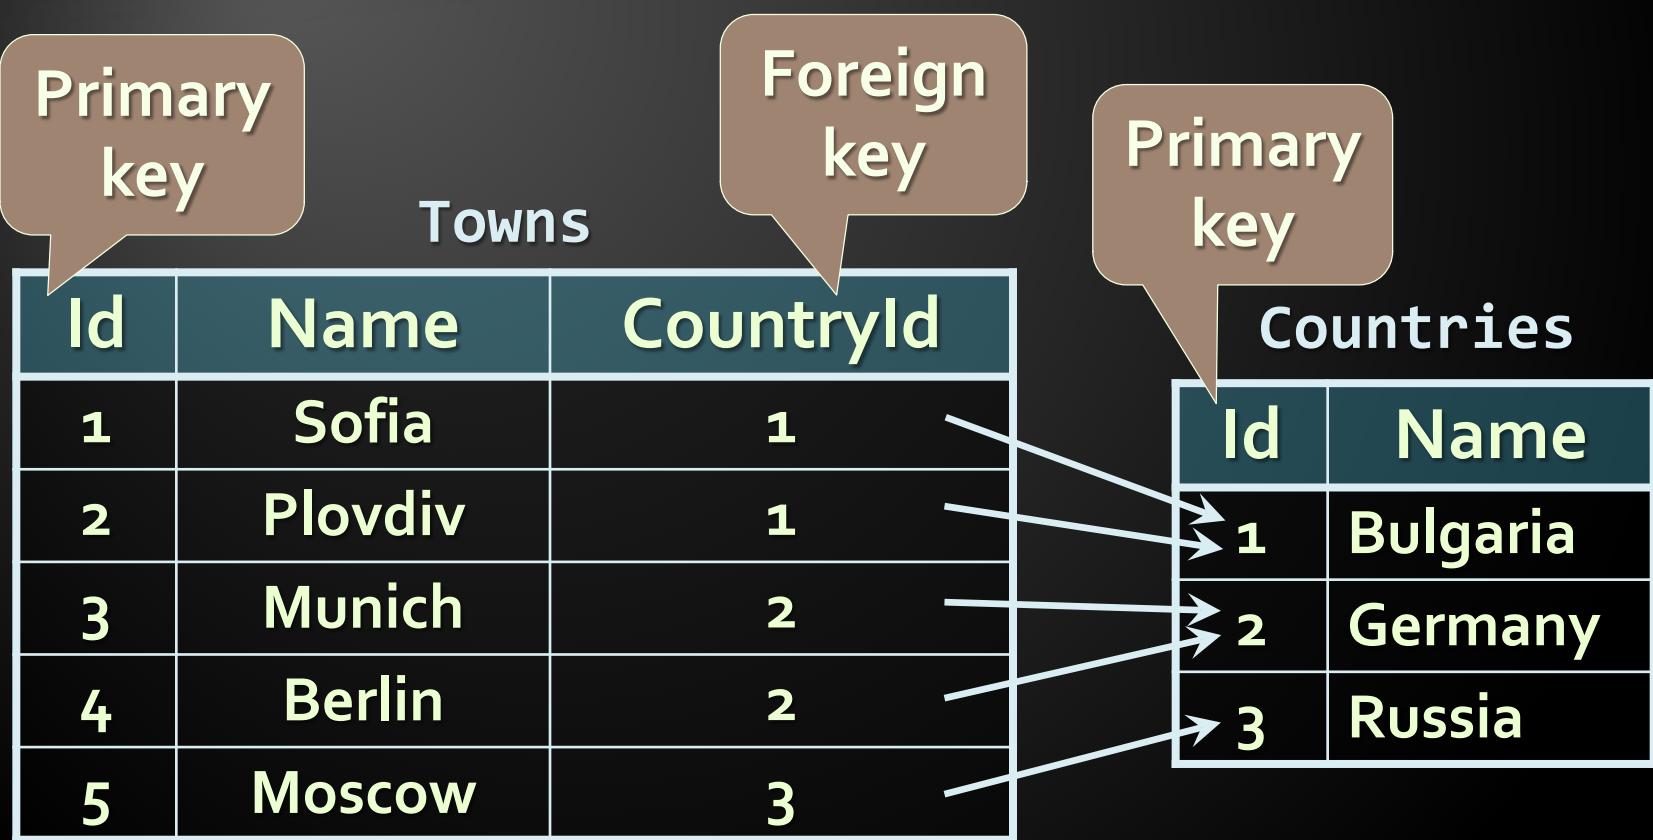
- ◆ Primary key is a column of the table that uniquely identifies its rows (usually it is a number)



Id	First Name	Last Name	Employer
1	Svetlin	Nakov	Telerik
2	Stephen	Forte	Telerik
3	Steve	Jobs	Apple

- ◆ Two records (rows) are different if and only if their primary keys are different
- ◆ The primary key can be composed by several columns (composite primary key)

- ◆ Relationships between tables are based on interconnections: primary key / foreign key



# Relationships (2)

- ◆ The foreign key is an identifier of a record located in another table (usually its primary key)
- ◆ By using relationships we avoid repeating data in the database
  - ◆ In the last example the name of the country is not repeated for each town (its number is used instead)
- ◆ Relationships have multiplicity:
  - ◆ One-to-many – e.g. country / towns
  - ◆ Many-to-many – e.g. student / course
  - ◆ One-to-one – e.g. example human / student

# Relationships' Multiplicity

- ◆ Relationship one-to-many (or many-to-one)
  - ◆ A single record in the first table has many corresponding records in the second table
  - ◆ Used very often

Towns

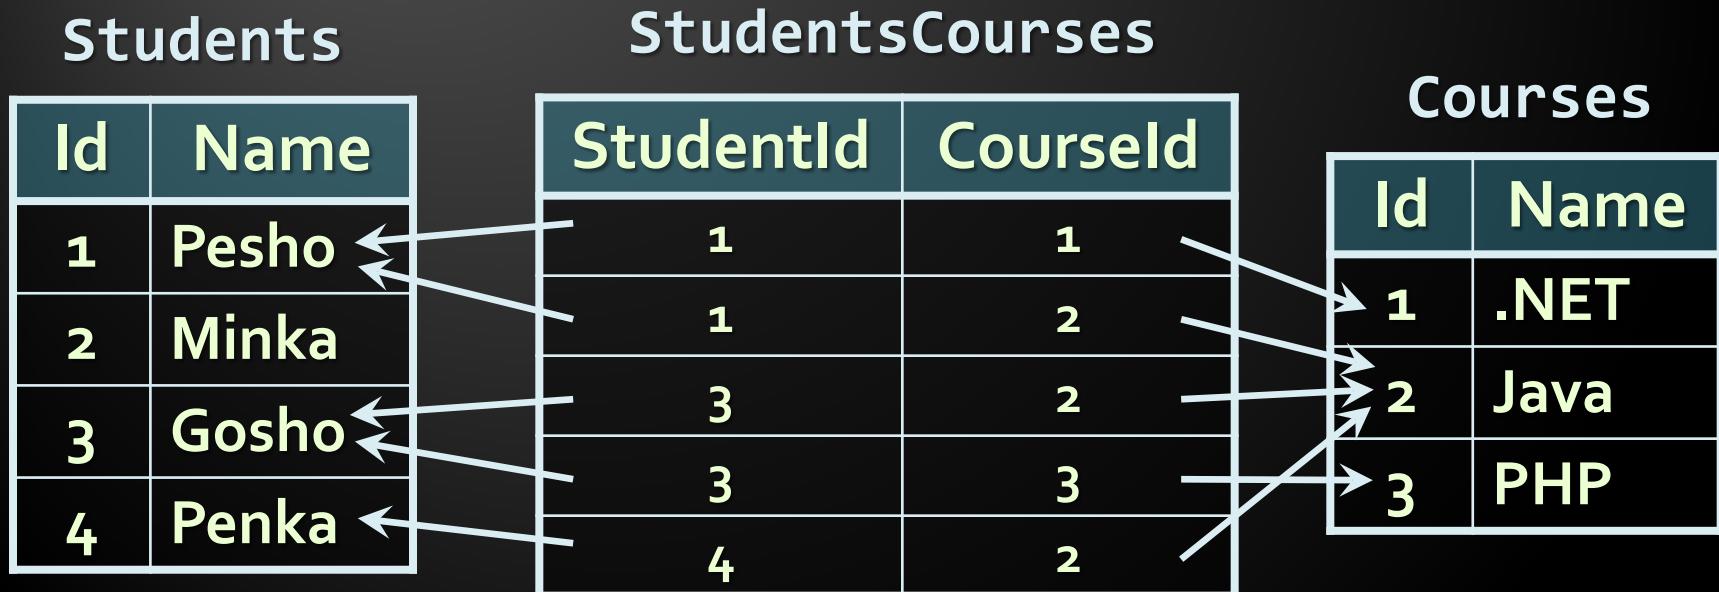
Id	Name	CountryId
1	Sofia	1
2	Plovdiv	1
3	Munich	2
4	Berlin	2
5	Moscow	3

Countries

Id	Name
1	Bulgaria
2	Germany
3	Russia

# Relationships' Multiplicity (2)

- ◆ Relationship many-to-many
  - ◆ Records in the first table have many corresponding records in the second one and vice versa
  - ◆ Implemented through additional table



# Relationships' Multiplicity (3)

## ◆ Relationship one-to-one

- A single record in a table corresponds to a single record in the other table
- Used to model inheritance between tables

Primary key

Primary & foreign  
key in the same time

Persons

Id	Name	Age
1	Ivan Daddy	72
2	Goiko Dude	26
3	Grand Mara	24

Professors

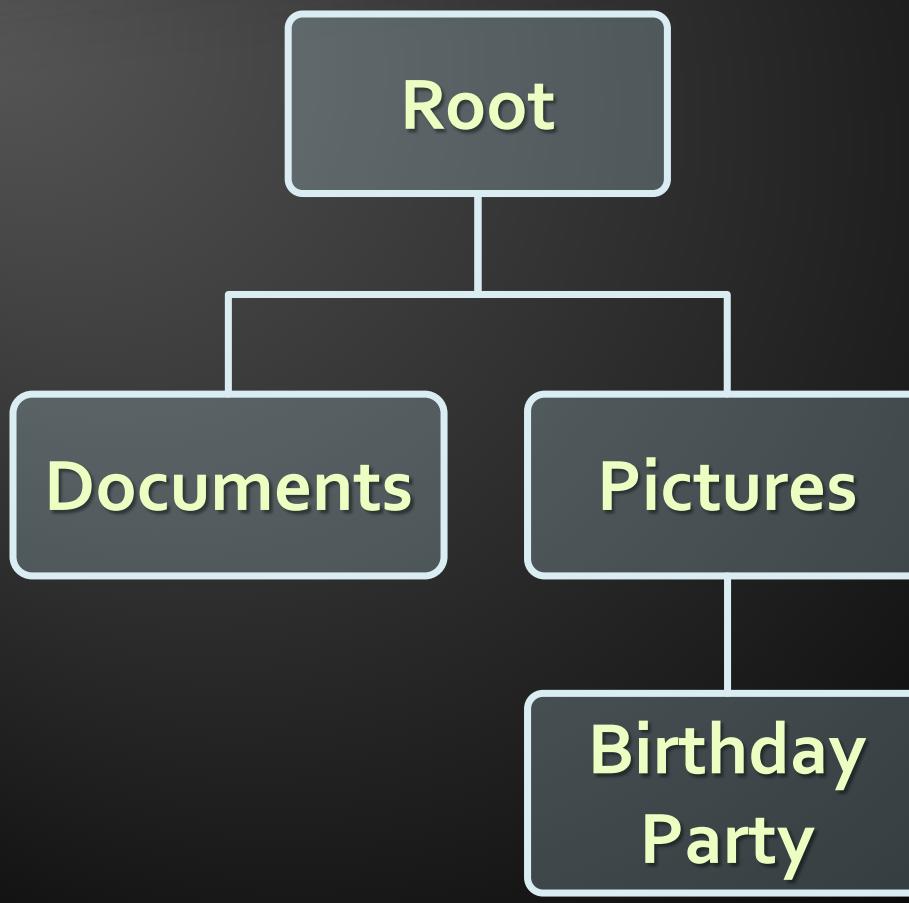
Id	Title
1	Ph.D.

Students

Id	Specialty
2	Computer Science
3	Chemistry

# Representing Hierarchical Data

- ◆ How do we represent trees and graphs?

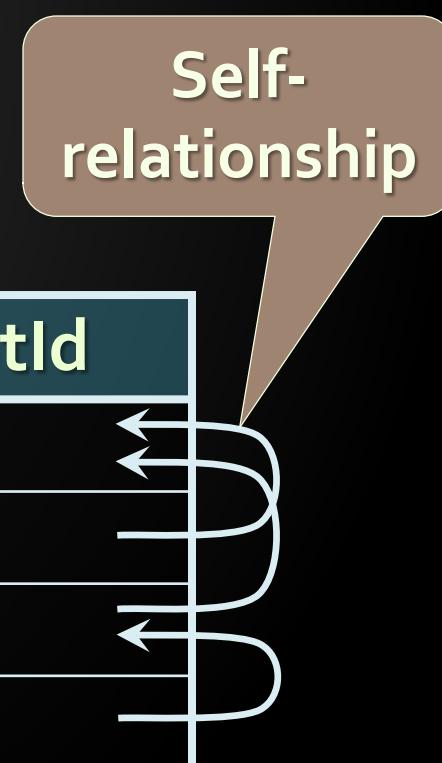


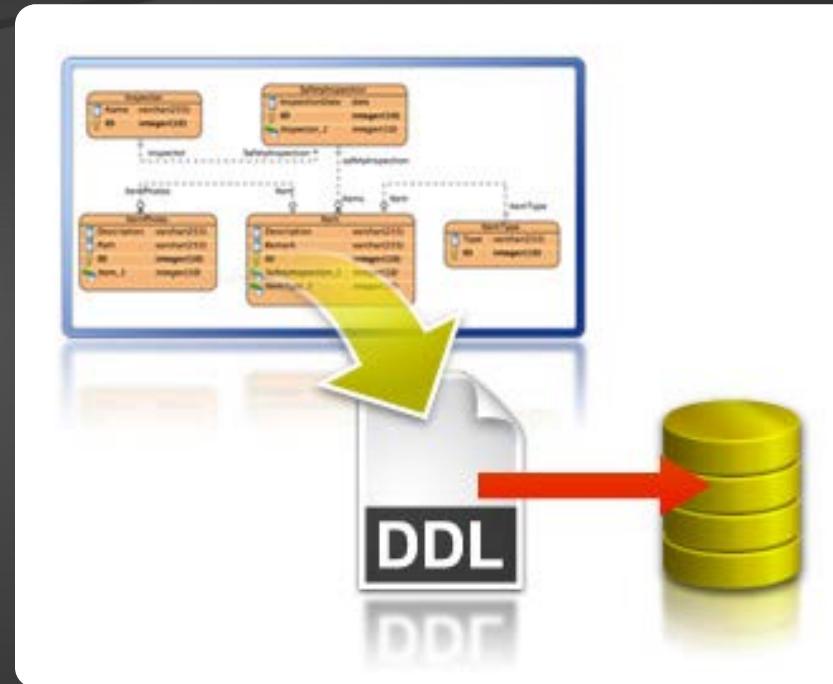
# Self-Relationships

- ◆ The primary / foreign key relationships can point to one and the same table
  - ◆ Example: employees in a company have a manager, who is also an employee

Primary key      Employees      Foreign key      Self-relationship

Id	Folder	ParentId
1	Root	(null)
2	Documents	1
3	Pictures	1
4	Birthday Party	3



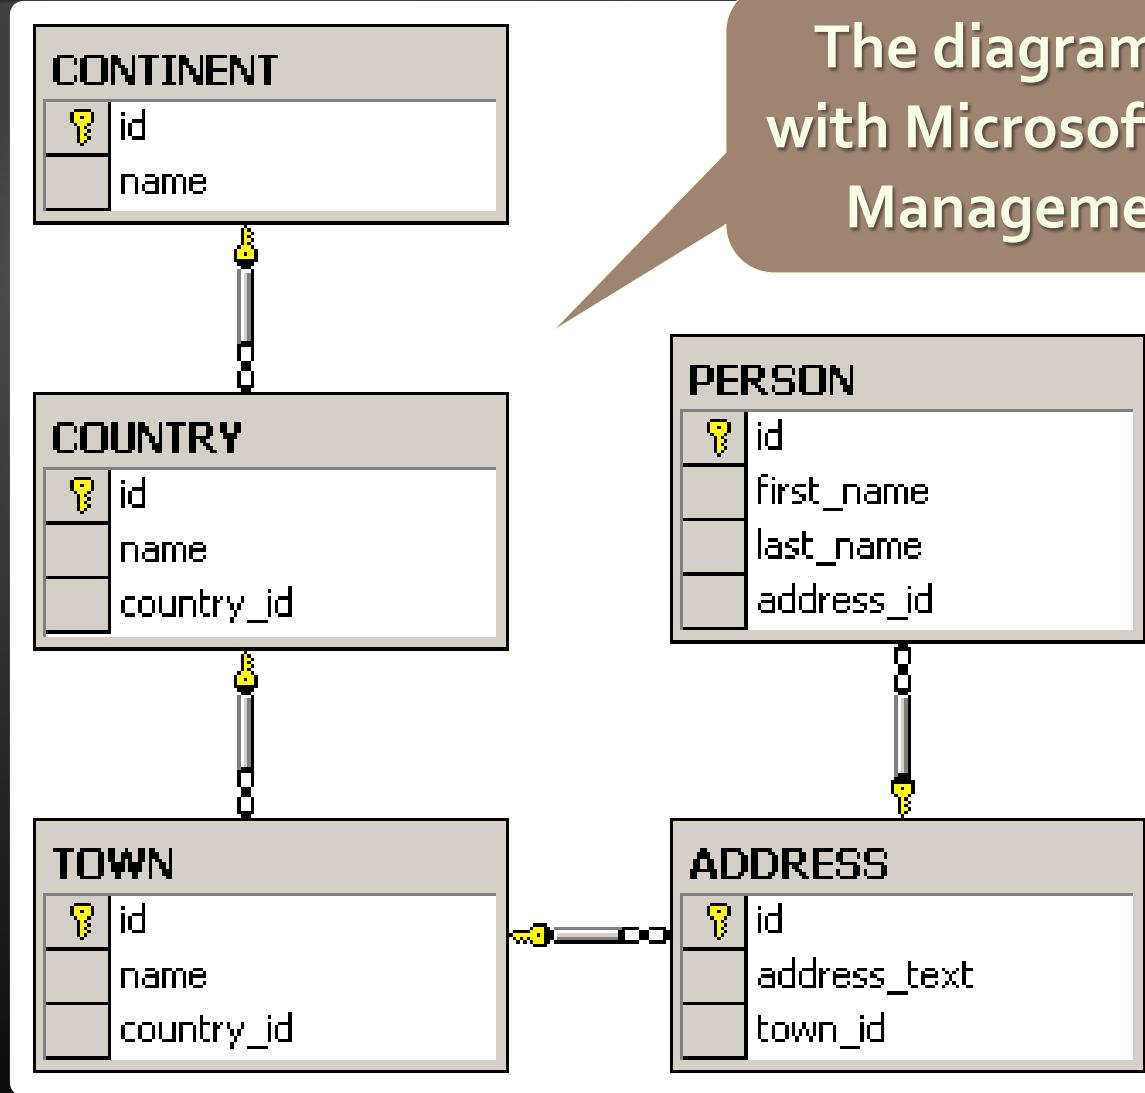


# E/R Diagrams

## Entity / Relationship Diagrams and DB Modeling Tools

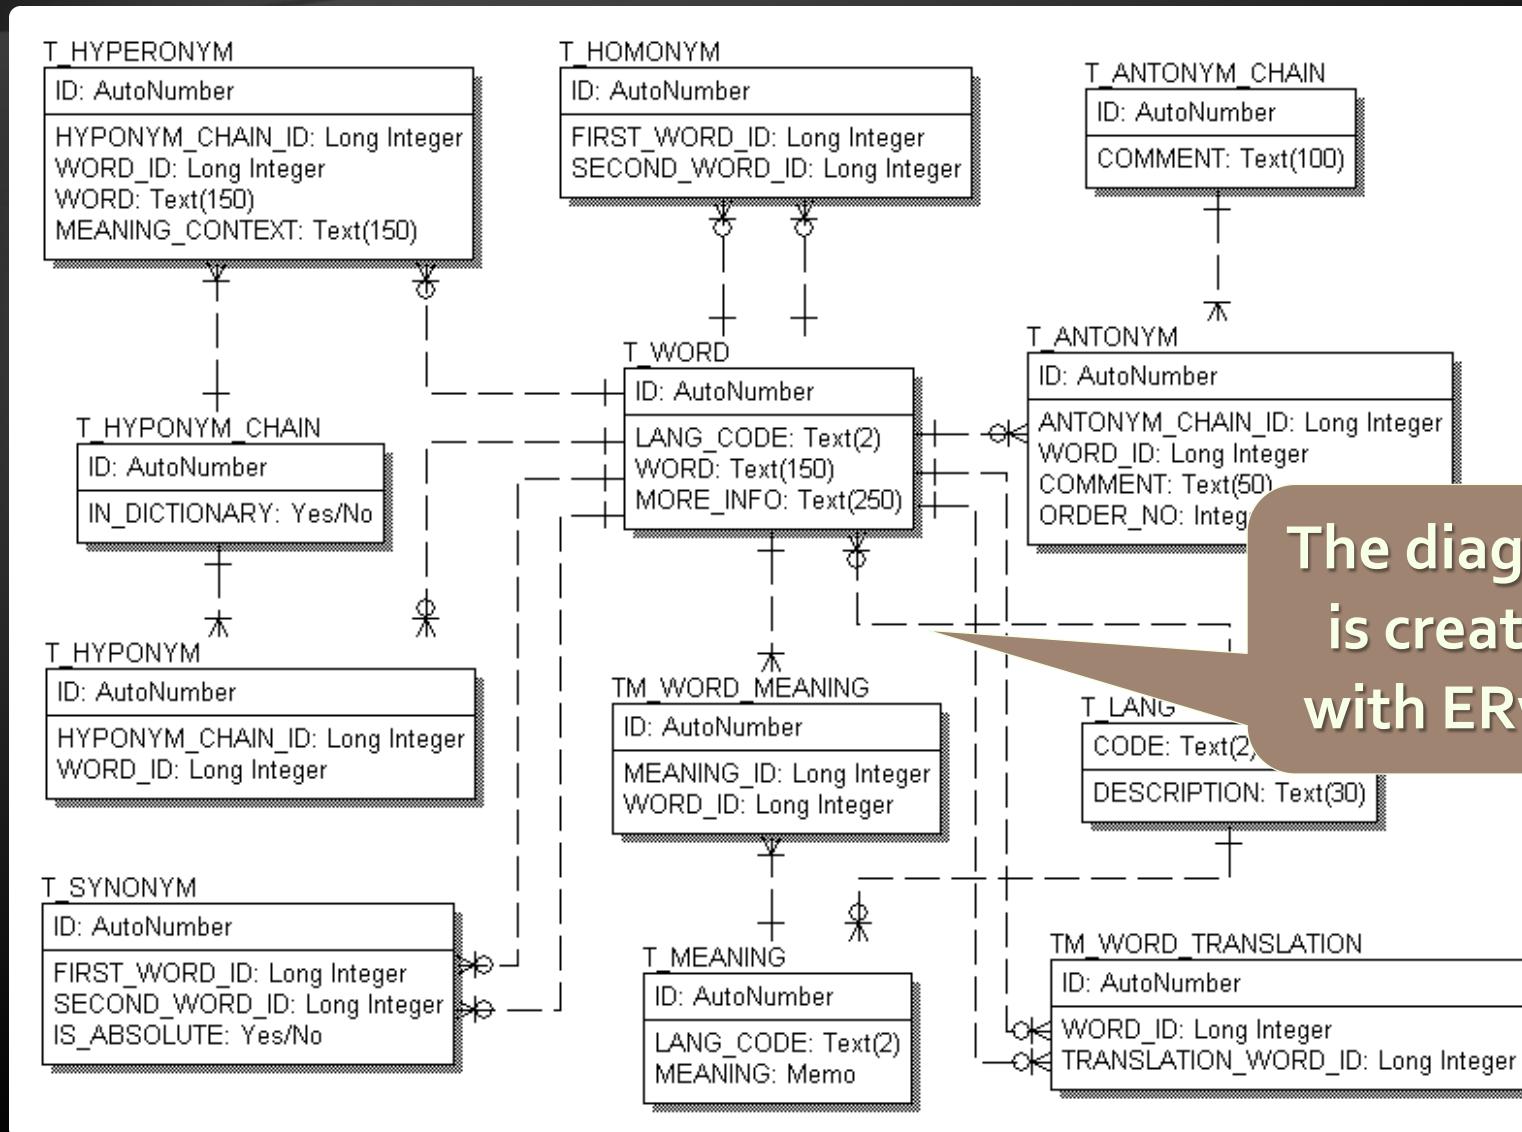
- ◆ Relational schema of a DB is the collection of:
  - The schemas of all tables
  - Relationships between the tables
  - Any other database objects (e.g. constraints)
- ◆ The relational schema describes the structure of the database
  - Doesn't contain data, but metadata
- ◆ Relational schemas are graphically displayed in Entity / Relationship diagrams (E/R Diagrams)

# E/R Diagrams – Examples



The diagram is created  
with Microsoft SQL Server  
Management Studio

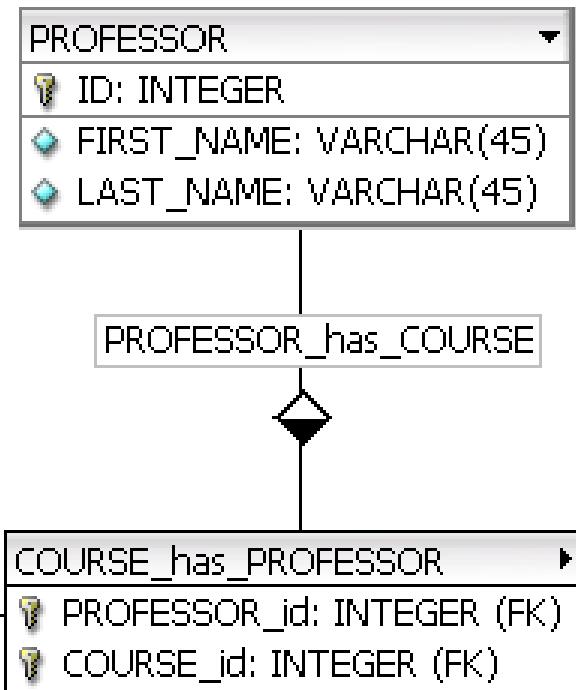
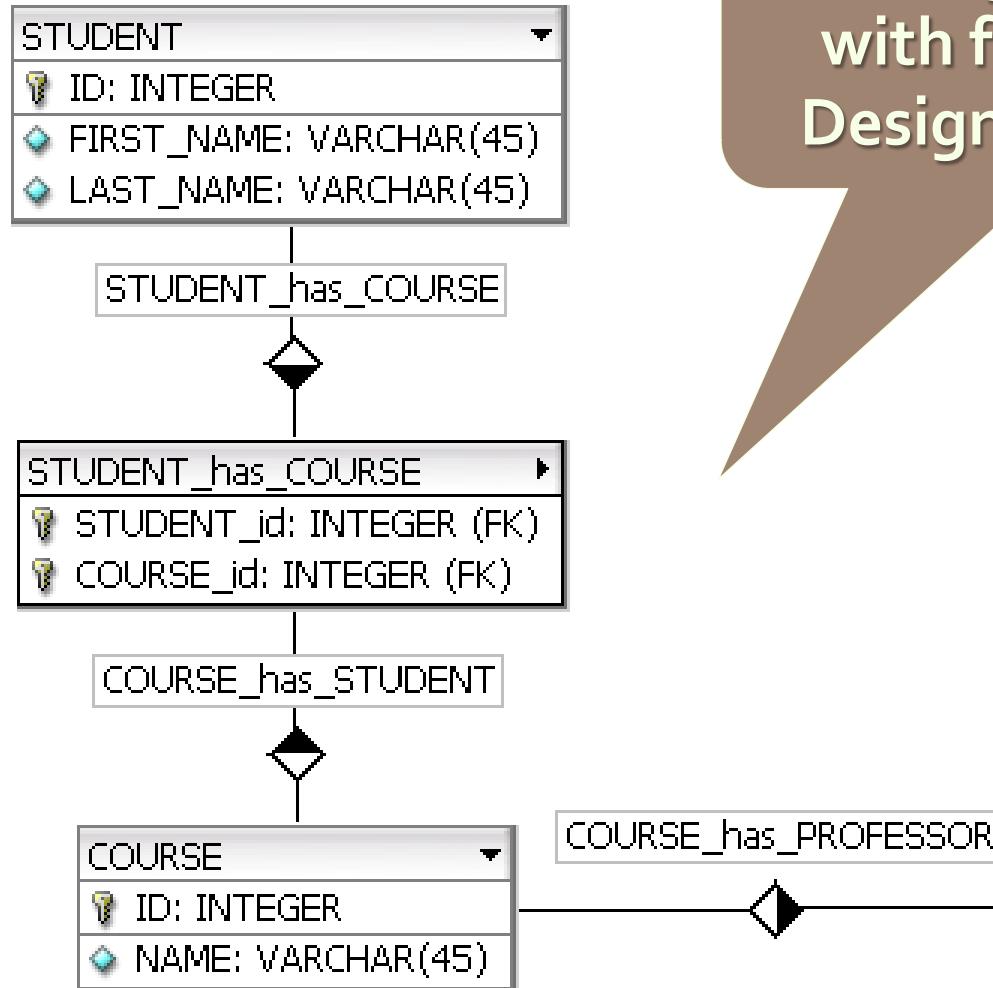
# E/R Diagrams – Examples (2)



The diagram  
is created  
with ERwin

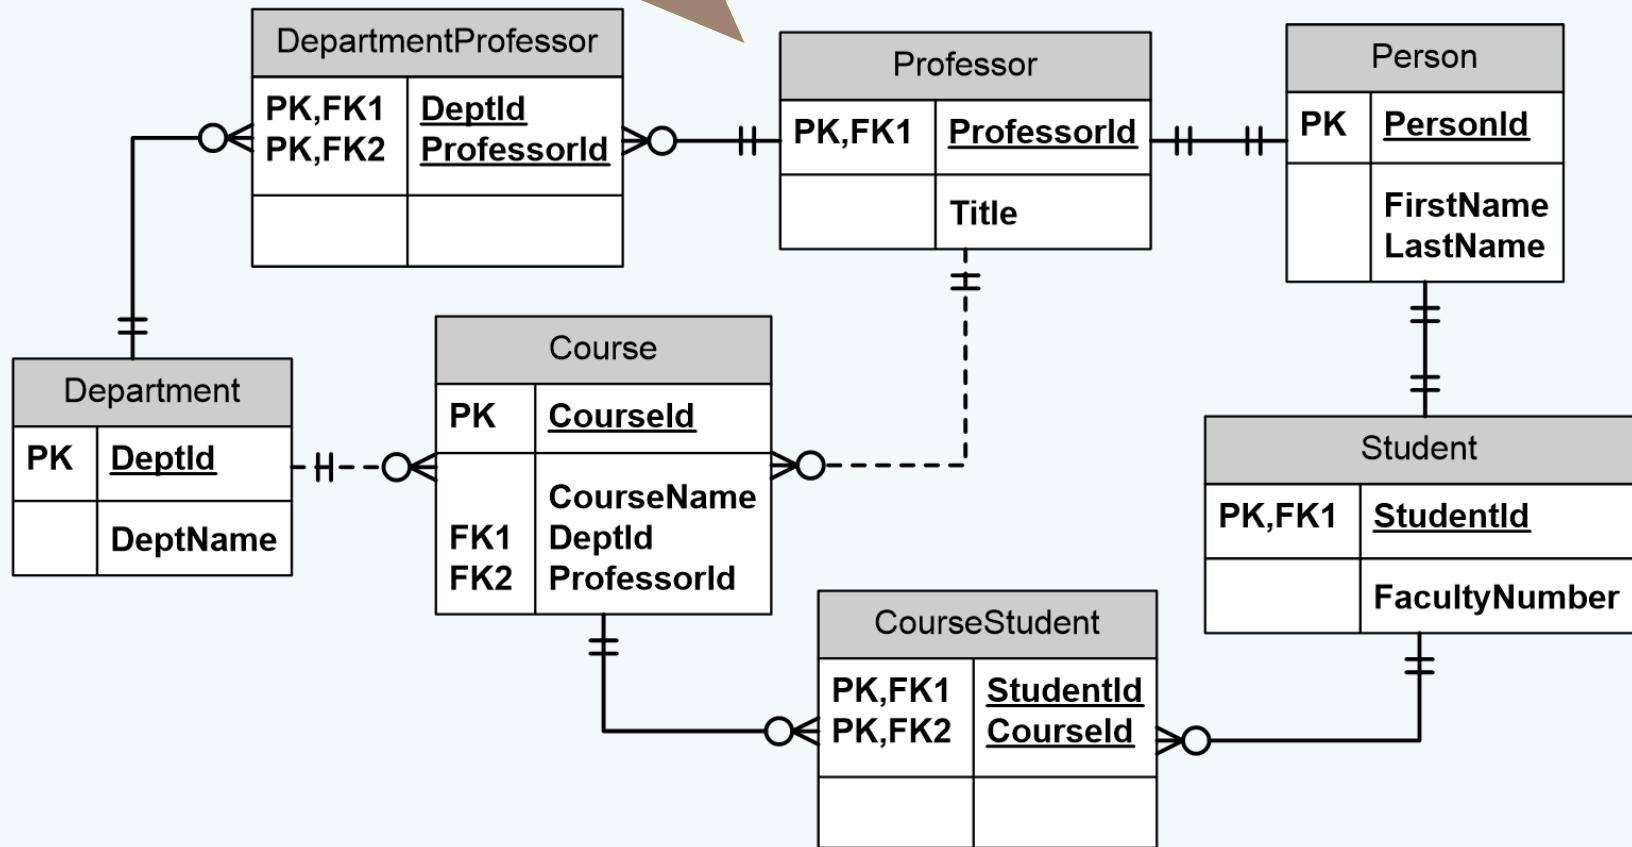
# E/R Diagrams – Examples (3)

The diagram is created  
with fabFORCE DB  
Designer for MySQL



# E/R Diagrams – Examples (4)

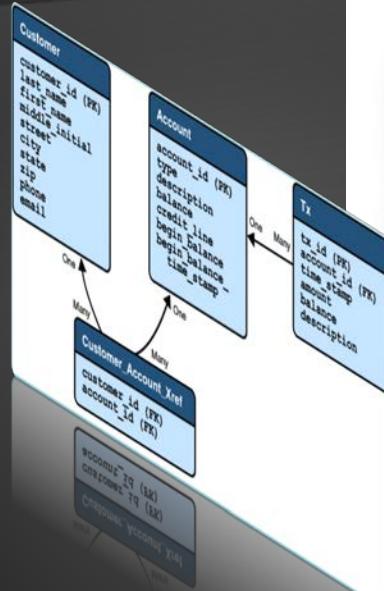
The diagram is  
created with MS Visio



# Tools for E/R Design

- ◆ Data modeling tools allow building E/R diagrams, generate / import DB schemas:
  - ◆ SQL Server Management Studio
  - ◆ MySQL Workbench
  - ◆ Oracle JDeveloper
  - ◆ Microsoft Visio
  - ◆ CASE Studio
  - ◆ Computer Associates ERwin
  - ◆ IBM Rational Rose





A	B	C	D
Affiliate	State	Members	Annual Fee
1 Norfolk	VA	205	50
2 Houston	TX	65	75
3 Manhattan	NY	657	60
4 Albany	NY	336	50
5 Washington	DC	453	50
6 Richmond	VA	432	25
7 Memphis	TN	77	70
8 Brooklyn	NY	578	65
9 Boston	MA	153	65
10 Waltham	MA	32	35
11 Schenectady	NY	43	85
12 Newark	NJ	235	75
13 Morristown	NJ	68	12
14		88	82
15		532	32
16		42	62
17		35	62
18		25	25
19		12	12
20		11	11
21		10	10
22		9	9
23		8	8
24		7	7
25		6	6
26		5	5
27		4	4
28		3	3
29		2	2
30		1	1
31		0	0

# DB Normalization

Avoiding Duplicated Data through  
Database Schema Normalization

# Normalization

- ◆ Normalization of the relational schema removes repeating data
- ◆ Non-normalized schemas can contain many data repetitions, e.g.

Product	Producer	Price	Category	Shop	Town
yoghurt	Mlexis Ltd.	0.67	food	store "Mente"	Sofia
bread "Dobrudja"	Bakery "Smoky"	0.85	food	store "Mente"	Sofia
beer "Zagorka"	Zagorka Corp.	0.68	soft drinks	stall "non-stop"	Varna
beer "Tuborg"	Shoumen Drinks Corp.	0.87	soft drinks	stall "non-stop"	Varna

# Normalization (6)

- ◆ Example of fully normalized schema (in 4<sup>th</sup> Normal Form):

Products

Id	Product	ProducerId	Price	CategoryId	ShopId	TownId
1	Youghurt	2	0.67	2	4	1
2	bread "Dobrudja"	3	0.55	2	4	1
3	rakia "Peshtera"	6	4.38	5	2	1
4	beer "Tuborg"	4	0.67	4	1	3

Producers

Id	Name
2	"Milk" Ltd.
4	"Zagorka" AD

Categories

Id	Name
4	beer
2	food

Shops

Id	Name
1	Billa
4	METRO

Towns

Id	Name
1	Sofia
3	Varna



# Other Database Objects

Constraints, Indices, SQL, Stored  
Procedures, Views, Triggers

# Integrity Constraints

- ◆ Integrity constraints ensure data integrity in the database tables
  - ◆ Enforce data rules which cannot be violated
- ◆ Primary key constraint
  - ◆ Ensures that the primary key of a table has unique value for each table row
- ◆ Unique key constraint
  - ◆ Ensures that all values in a certain column (or a group of columns) are unique

# Integrity Constraints (2)

- ◆ Foreign key constraint
  - ◆ Ensures that the value in given column is a key from another table
- ◆ Check constraint
  - ◆ Ensures that values in a certain column meet some predefined condition
  - ◆ Examples:

```
(hour >= 0) AND (hour < 24)
```

```
name = UPPER(name)
```

- ◆ Indices speed up searching of values in a certain column or group of columns
  - ◆ Usually implemented as B-trees
- ◆ Indices can be built-in the table (clustered) or stored externally (non-clustered)
- ◆ Adding and deleting records in indexed tables is slower!
  - ◆ Indices should be used for big tables only (e.g. 50 000 rows)

# The SQL Language

- ◆ **SQL (Structured Query Language)**
  - ◆ Standardized declarative language for manipulation of relational databases
  - ◆ SQL-99 is currently in use in most databases
  - ◆ <http://en.wikipedia.org/wiki/SQL#Standardization>
- ◆ **SQL language supports:**
  - ◆ Creating, altering, deleting tables and other objects in the database
  - ◆ Searching, retrieving, inserting, modifying and deleting table data (rows)

# The SQL Language (2)

- ◆ SQL consists of:
  - ◆ DDL – Data Definition Language
    - ◆ CREATE, ALTER, DROP commands
  - ◆ DML – Data Manipulation Language
    - ◆ SELECT, INSERT, UPDATE, DELETE commands
- ◆ Example of SQL SELECT query:

```
SELECT Towns.Name, Countries.Name  
FROM Towns, Countries  
WHERE Towns.CountryId = Countries.Id
```

- ◆ **Stored procedures (database-level procedures)**
  - ◆ Consist of SQL-like code stored in the database
  - ◆ Code executed inside the database server
  - ◆ Much faster than an external code
  - ◆ Data is locally accessible
  - ◆ Can accept parameters
  - ◆ Can return results
    - ◆ Single value
    - ◆ Record sets



# Stored Procedures (2)

- ◆ Stored procedures are written in a language extension of SQL
  - ◆ T-SQL – in Microsoft SQL Server
  - ◆ PL/SQL – in Oracle
- ◆ Example of stored procedure in Oracle PL/SQL:

```
CREATE OR REPLACE PROCEDURE
spInsertCountry(countryName varchar2) IS
BEGIN
    INSERT INTO Countries(Name)
    VALUES(countryName);
END;
```

- ◆ Views are named SQL SELECT queries which are used as tables
  - ◆ Simplify data access
  - ◆ Facilitate writing of complex SQL queries
- ◆ Used also to apply security restrictions:
  - ◆ E.g. a certain user isn't given permissions on any of the tables in the database
  - ◆ The user is given permissions on few views (subset of DB) and few stored procedures only

# Views – Example

## Companies

<b>Id</b>	<b>Company</b>	<b>TownId</b>
<b>1</b>	Mente LTD	<b>1</b>
<b>2</b>	BulkSoft Inc.	<b>2</b>
<b>3</b>	HardSoft Corp.	<b>4</b>
<b>4</b>	Sputnik Corp.	<b>3</b>

## Towns

<b>Id</b>	<b>Town</b>	<b>CountryId</b>
<b>1</b>	Sofia	<b>1</b>
<b>2</b>	New York	<b>3</b>
<b>3</b>	Moscow	<b>2</b>
<b>4</b>	Plovdiv	<b>1</b>

## Countries

<b>Id</b>	<b>Country</b>
<b>1</b>	Bulgaria
<b>2</b>	Russia
<b>3</b>	USA

# Views – Example (2)

```
CREATE VIEW V_BGCompanies AS
SELECT
    Companies.Id AS Id,
    Companies.Company AS Company
FROM Companies INNER JOIN
    (Towns INNER JOIN Countries ON
        Towns.CountryId = Countries.Id)
    ON Companies.TownId = Towns.Id
WHERE
    Countries.Country = "Bulgaria";
```



V\_BGCompanies

Id	Company
1	Mente Ltd.
3	HardSoft Corp.

- ◆ Triggers are special stored procedures that are activate when some event occurs, for instance:
  - When inserting a record
  - When changing a record
  - When deleting a record
- ◆ Triggers can perform additional data processing of the affected rows, e.g.
  - To change the newly added data
  - To maintain logs and history



# Triggers – Example

- ◆ We have a table holding company names:

```
CREATE TABLE Companies(  
    Id number NOT NULL,  
    Name varchar(50) NOT NULL)
```

- ◆ A trigger that appends "Ltd." at the end of the name of a new company:

```
CREATE OR REPLACE TRIGGER trg_Companies_INSERT  
    BEFORE INSERT ON Company  
    FOR EACH ROW  
BEGIN  
    :NEW.Name := :NEW.Name || ' Ltd.';  
END;
```

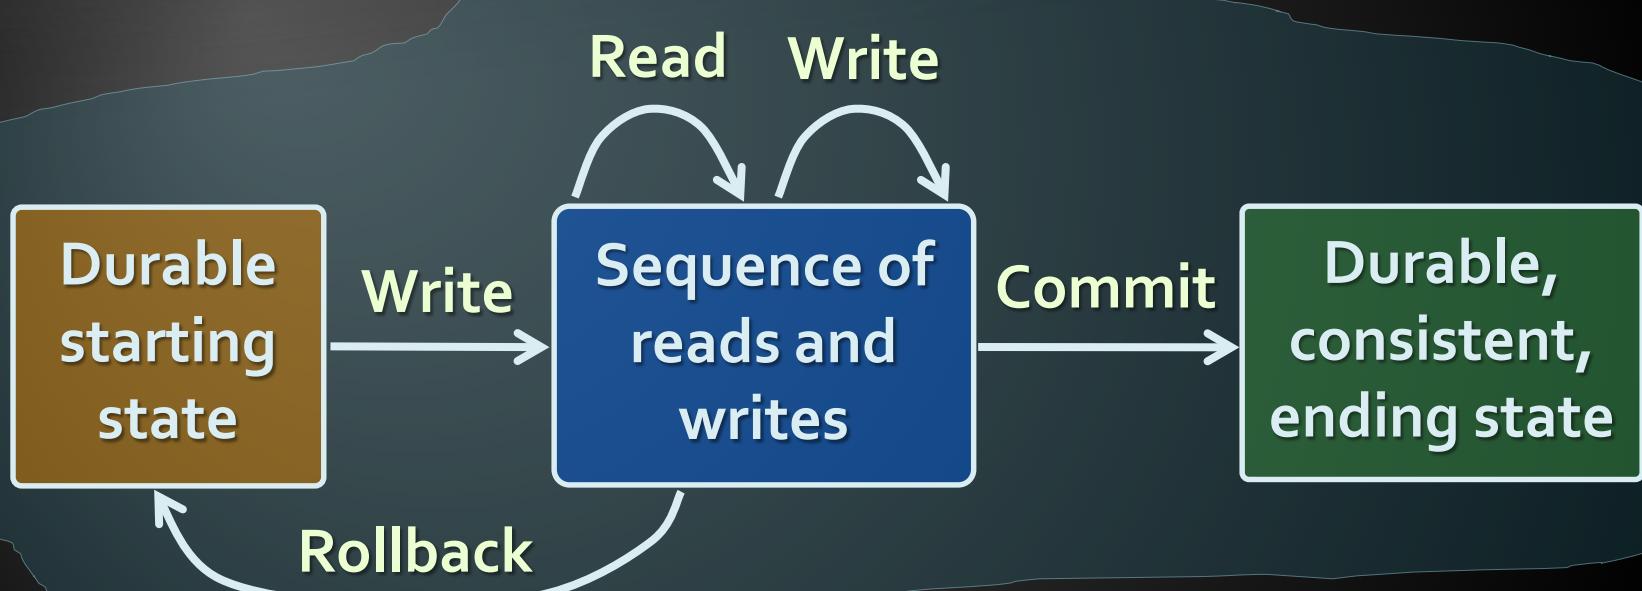
# Transactions

ACID Transactions and Isolation



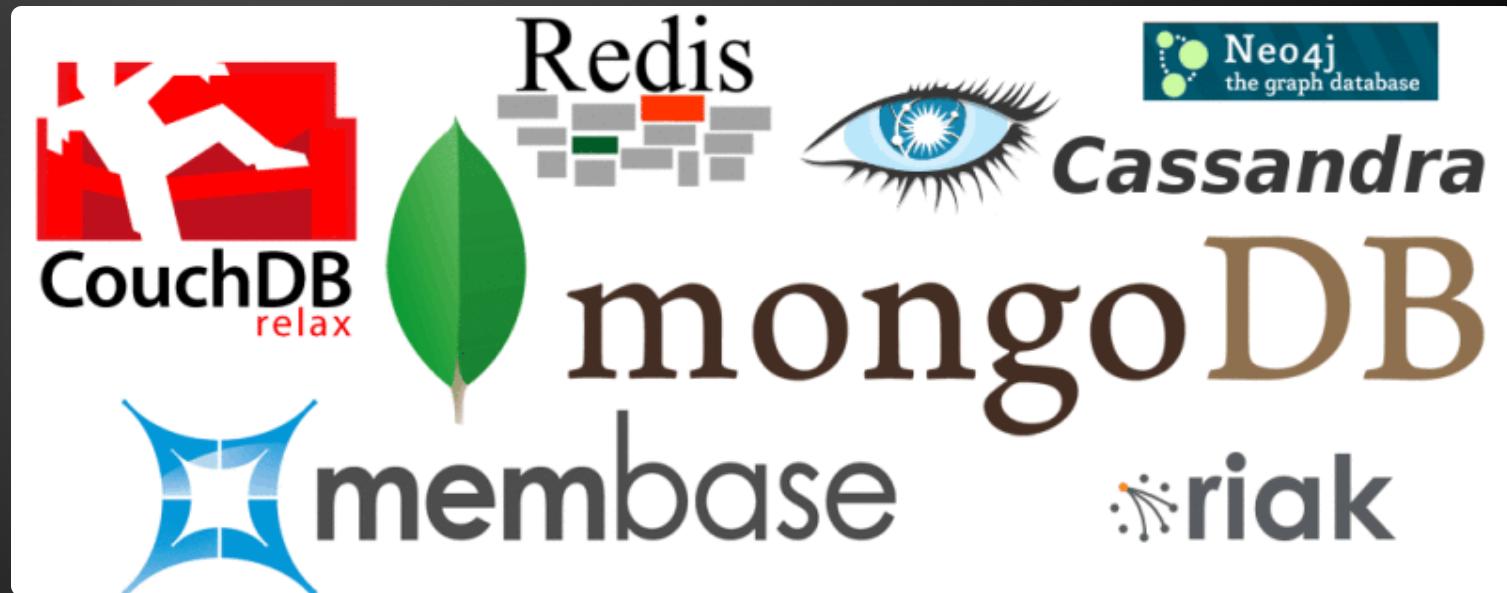
- ◆ Transactions are a sequence of database operations which are executed as a single unit:
  - Either all of them execute successfully
  - Or none of them is executed at all
- ◆ Example:
  - A bank transfer from one account into another (withdrawal + deposit)
  - If either the withdrawal or the deposit fails the entire operation should be cancelled

# DB Transactions Lifecycle



# Transactions Behavior

- ◆ Transactions guarantee the consistency and the integrity of the database
  - All changes in a transaction are temporary
  - Changes become final when COMMIT is successfully executed
  - At any time all changes done in the transaction can be canceled by executing ROLLBACK
- ◆ All operations are executed as a single unit
  - Either all of them pass or none of them

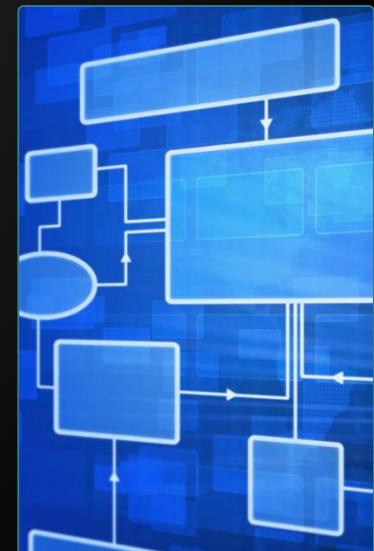
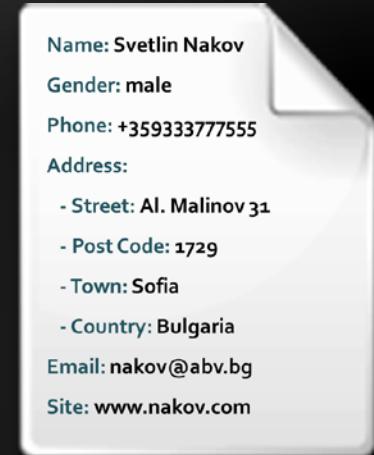


# NoSQL Databases

Non-Relational Database Systems

# Non-Relational Data Models

- ◆ Document model
  - ◆ Set of documents, e.g. JSON strings
- ◆ Key-value model
  - ◆ Set of key-value pairs
- ◆ Hierarchical key-value
  - ◆ Hierarchy of key-value pairs
- ◆ Wide-column model
  - ◆ Key-value model with schema
- ◆ Object model
  - ◆ Set of OOP-style objects



# What is NoSQL Database?

- ◆ NoSQL (non-relational) databases
  - ◆ Use document-based model (non-relational)
  - ◆ Schema-free document storage
    - ◆ Still support CRUD operations (create, read, update, delete)
    - ◆ Still support indexing and querying
    - ◆ Still supports concurrency and transactions
  - ◆ Highly optimized for append / retrieve
  - ◆ Great performance and scalability
  - ◆ NoSQL == “No SQL” or “Not Only SQL”?



## ◆ Relational databases

- ◆ Data stored as table rows
- ◆ Relationships between related rows
- ◆ Single entity spans multiple tables
- ◆ RDBMS systems are very mature, rock solid



## ◆ NoSQL databases

- ◆ Data stored as documents
- ◆ Single entity (document) is a single record
- ◆ Documents do not have a fixed structure



# Relational vs. NoSQL Models

## Relational Model

Name	Svetlin Nakov
Gender	male
Phone	+359333777555
Email	nakov@abv.bg
Site	www.nakov.com

\*  
1

Street	Al. Malinov 31
Post Code	1729

\*  
1

Town	Sofia
------	-------

\*  
1

Country	Bulgaria
---------	----------

## Document Model

Name: Svetlin Nakov

Gender: male

Phone: +359333777555

Address:

- Street: Al. Malinov 31

- Post Code: 1729

- Town: Sofia

- Country: Bulgaria

Email: nakov@abv.bg

Site: www.nakov.com

# NoSQL Database Systems

- ◆ Redis



- ◆ Ultra-fast in-memory data structures server

- ◆ MongoDB



- ◆ Mature and powerful JSON-document database

- ◆ CouchDB



- ◆ JSON-based document database with REST API

- ◆ Cassandra



- ◆ Distributed wide-column database

- ◆ DB Ranking: <http://db-engines.com/en/ranking>

## Questions?

1. What database models do you know?
2. Which are the main functions performed by a Relational Database Management System (RDBMS)?
3. Define what is "table" in database terms.
4. Explain the difference between a primary and a foreign key.
5. Explain the different kinds of relationships between tables in relational databases.
6. When is a certain database schema normalized? What are the advantages of normalized databases?

7. What are database integrity constraints and when are they used?
8. Point out the pros and cons of using indexes in a database.
9. What's the main purpose of the SQL language?
10. What are transactions used for? Give an example.
11. What is a NoSQL database?
12. Explain the classical non-relational data models.
13. Give few examples of NoSQL databases and their pros and cons.

# Free Trainings @ Telerik Academy

- ◆ C# Programming @ Telerik Academy

- ◆ [csharpfundamentals.telerik.com](http://csharpfundamentals.telerik.com)



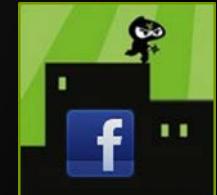
- ◆ Telerik Software Academy

- ◆ [academy.telerik.com](http://academy.telerik.com)



- ◆ Telerik Academy @ Facebook

- ◆ [facebook.com/TelerikAcademy](https://facebook.com/TelerikAcademy)



- ◆ Telerik Software Academy Forums

- ◆ [forums.academy.telerik.com](http://forums.academy.telerik.com)

